



Image Processing & Vision

Lecture 05: Feature Descriptors

Hak Gu Kim

hakgukim@cau.ac.kr

Immersive Reality & Intelligent Systems Lab (IRIS LAB)

Graduate School of Advanced Imaging Science, Multimedia & Film (GSAIM)

Chung-Ang University (CAU)

03 Apr. 2023

Recap: Local Feature

- Global template matching → **Local feature detection**
- Consider the problem of finding images of an elephant using a template
 - An elephant looks different from different viewpoints
 - What happens if parts of an elephant are obscured from view by trees, rocks, other elephants?



Template




Image




Normalized correlation map

Recap: Edges

- An edge is a place of rapid change in the image intensity function




Image



Intensity function
along horizontal scanline



*Edges corresponding to
extreme of derivatives*




1st derivative
in horizontal axis

Recap: DoG & LoG


- **Two Generic Edge Detection Approaches:**
 - Local extrema of a first derivative operator
 - Zero crossings of a second derivative operator
- **Laplacian of Gaussian** is based on a zero crossings of a second derivative operator method



G_{σ}



∇G_{σ}




$\nabla^2 G_{\sigma}$

Recap: Canny Edge Detection


- ① Apply **directional derivatives of Gaussian**
- ② Compute **gradient magnitude** and **gradient direction**
- ③ **Non-maximum suppression (NMS)**
- ④ **Link and threshold**



Original image



Gradient magnitude




Result of non-maximum suppression

Recap: Corner


- **Key Property:**
 - In the region around a corner, image gradient has two or more dominant directions
 - Corners are robust and distinctive

Flat




No changes
in all directions

Edge



No changes
along the edge direction

Corner



Significant changes
in all directions



Recap: Harris Corner Detection

- ① Compute **image gradient** over small region
- ② Compute the **covariance matrix**
- ③ Compute **eigenvectors** and **eigenvalues**
- ④ Use **threshold on eigenvalues** to detect corners

Sum over local window
region around corner

Gradient with respect to
each direction (x or y)

$$\mathbf{M} = \begin{bmatrix} \sum_{p \in P} I_x I_x & \sum_{p \in P} I_x I_y \\ \sum_{p \in P} I_y I_x & \sum_{p \in P} I_y I_y \end{bmatrix}$$




Recap: Harris Corner Detection

- **Visualization** of Eigenvalues and Eigenvectors as an **Ellipse**
 - **Eigenvectors** determines the **orientation of the ellipse**
 - **Eigenvalues** determines the axis **lengths of the ellipse**
 - $\frac{1}{\sqrt{\lambda_{max}}}$: The direction of the fastest change
 - $\frac{1}{\sqrt{\lambda_{min}}}$: The direction of the slowest change

$$\mathbf{M} = \begin{bmatrix} \sum_{p \in P} I_x I_x & \sum_{p \in P} I_x I_y \\ \sum_{p \in P} I_y I_x & \sum_{p \in P} I_y I_y \end{bmatrix} = \begin{bmatrix} \mathbf{v}_1 & \mathbf{v}_2 \end{bmatrix} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} \begin{bmatrix} \mathbf{v}_1 & \mathbf{v}_2 \end{bmatrix}^{-1}$$

Eigenvalues Eigenvectors



Recap: Harris Corner Detection


- Use **threshold on eigenvalues** to detect corners

- **Cornerness:**


$$\begin{aligned} C &= \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2 \\ &= \det(\mathbf{M}) - \alpha \cdot \text{tr}(\mathbf{M})^2 \end{aligned}$$

$$0.04 \leq \alpha \leq 0.06$$


- Flat region: $|C| \leq \epsilon$
- Edge: $C < 0$
- Corner: $C > 0$



Recap: Harris Corner Detection



Corner response map



Corner detection result ($\sigma = 1$)



Topics

- Feature Descriptors
 - Basic Feature Descriptors
 - SIFT

***Note:** Many of these slides in this course were adapted from Computer Vision at CMU (16-385) and UBC (CPSC425)



Difficulty of Matching: Photometric Transformations

- How do we match the correspondences between two images?




Difficulty of Matching: Geometric Transformations

- How do we match the correspondences between two images?



Good Local Features

- How do we know which **corner** goes with which?



Good Local Features

- Patch around the local feature is much more **informative**






Image Intensity

- Just use the **pixel values** of the patch




- **Pros:** Perfectly fine if geometry and appearance is unchanged
- **Cons:** Very sensitive to absolute intensity values

Image Gradients & Edges

- Just use the **Edge values/signs** of the patch



Local patch





Vector of edge values/signs

- **Pros:** Invariant to absolute intensity values
- **Cons:** Very sensitive to deformations

Color Histogram




- Count the colors in the image using a histogram



- **Pros:** Invariant to changes in scale and rotation
- **Cons:** Ignoring its shape and texture (i.e., spatial layout)

Color Histogram


- Count the colors in the image using a histogram



- **Pros:** Invariant to changes in scale and rotation
- **Cons:** Ignoring its shape and texture (i.e., spatial layout)

Spatial Histograms

- Count the colors in **each local region** using a histogram



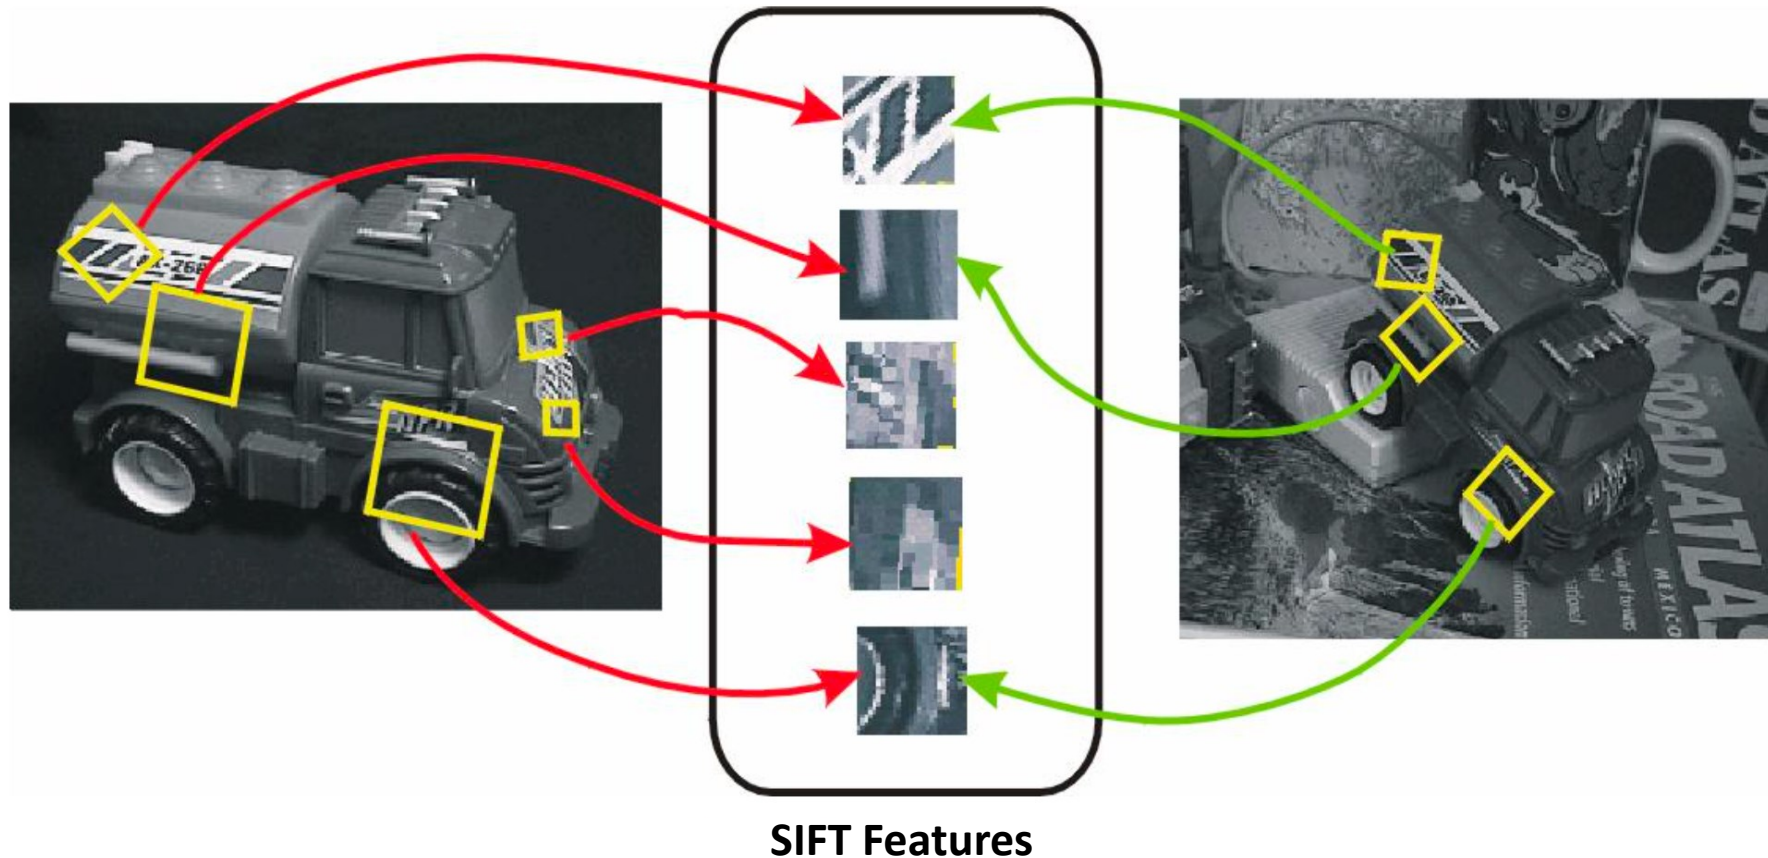
- **Pros:** Invariant to some deformations and retains rough spatial layout
- **Cons:** Sensitive to rotation

Summary of Features

Representation	Result	Approach	Technique
Intensity	Dense (2D)	Template Matching	Normalized correlation, Sum of squared difference (SSD)
Edge	Relatively sparse (1D)	Derivatives	$\nabla^2 G$, Canny edge
Corner / Blob	Sparse (0D)	Locally distinct features	Harris corner, SIFT

Invariant Local Features

- Image content is transformed into local feature coordinates that are **invariant to translation, rotation, scale, and other imaging parameters**




Advantages of Invariant Local Features

- **Locality:** Features are local, so robust to occlusion and clutter
- **Distinctiveness:** Individual features can be matched to a large-scale database of objects
- **Quantity:** Many features can be generated for even small objects
- **Efficiency:** Close to real-time performance

Scale Invariant Feature Transform (SIFT)


- The scale-invariant feature transform (SIFT) is an algorithm to **detect**, **describe**, and **match** local features in images
- SIFT describes both a **detector** and **descriptor**
 - Applications: object recognition, robotic mapping and navigation, image stitching, 3D modeling, gesture recognition, video tracking, *etc.*

- ① Multi-scale extrema detection
- ② Keypoint localization
- ③ Orientation assignment
- ④ Keypoint descriptor




SIFT: Multi-Scale Extrema Detection

- Begin by detecting points of interest (i.e., keypoints)
 - The image is convolved with Gaussian filters at different scales




SIFT: Multi-Scale Extrema Detection

- Begin by detecting points of interest (i.e., keypoints)
 - The image is convolved with Gaussian filters at different scales
 - Then the difference of successive Gaussian-blurred images are taken



SIFT: Multi-Scale Extrema Detection

- Begin by detecting points of interest (i.e., keypoints)
 - The image is convolved with Gaussian filters at different scales
 - Then the difference of successive Gaussian-blurred images are taken
 - Keypoints are taken as maxima/minima of the DoG at multiple scales



SIFT: Keypoint Localization

- Scale-space extrema detection produces too many keypoint candidates, some of which are **unstable**
- In keypoint localization, we reject points which are **low contrast** (and are therefore **sensitive to noise**) or **poorly localized** along an edge
- How do we decide whether a keypoint is poorly localized or well-localized?
 - In SIFT, compute the ratio of the eigenvalues of covariance matrix and check if it is greater than a threshold



SIFT: Keypoint Localization



(a) 233 x 189 image

(b) 832 DoG extrema in (a)

(c) 729 left after peak value threshold



(d) 536 left after testing ratio of principal curvatures

SIFT: Orientation Assignment

- Each keypoint is assigned one or more orientations based on local image gradient directions
- This is the key step in achieving invariance to rotation on the Gaussian-smoothed image, L



$$m(x, y) = \sqrt{(L(x + 1, y) - L(x - 1, y))^2 + (L(x, y + 1) - L(x, y - 1))^2}$$

$$\theta(x, y) = \tan^{-1}[(L(x + 1, y) - L(x - 1, y)) / (L(x, y + 1) - L(x, y - 1))]$$





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- **Gradient orientation:** Direction of arrows
- **Gradient magnitude:** Length of arrows





SIFT: Orientation Assignment

- Gradient magnitude by a Gaussian kernel



SIFT: Orientation Assignment

- Orientation Assignment method:
 - Create **histogram of local gradient directions** computed at selected scale
 - Histogram of 36 bins (10 degree increments)
 - Assign **canonical orientation** at peak of smoothed histogram
 - Gaussian-weighted voting
 - Each key specifies stable 2D coordinates: $(x, y, scale, orientation)$
 - Highest peak and peaks above 80% of highest also considered for calculating dominant orientations




SIFT: Keypoint Descriptor

- **Keypoint detection**
 - How to assign a location, scale and orientation to each keypoint
 - This ensured invariance to image location, scale and rotation
- **Keypoint description**
 - The descriptor is highly distinctive and partially invariant to the remaining variations such as illumination, 3D viewpoint, *etc.*

SIFT: Keypoint Descriptor

- Thresholded image gradients are sampled over 16×16 array of locations in scale space (weighted by a Gaussian with sigma half the size of the window)
- Create array of orientation histograms
- 8 Orientations x (4 x 4) histogram array



SIFT: Keypoint Descriptor

- Descriptor is normalized to unit length (i.e., magnitude of 1) to reduce the effects of illumination change
 - If brightness values are scaled (multiplied) by a constant, the gradients are scaled by the same constant, and the normalization cancels the change
 - If brightness values are increased/decreased by a constant, the gradients do not change

Summary: SIFT

① Scale-space representation and local extreme detection

- Use DoG/LoG Pyramid
- 3 scales/octave, down-sample by factor of 2 each octave

② Keypoint localization

- Select stable keypoints by thresholding on the magnitude of extrema and ratio of principal curvatures

③ Keypoint orientation assignment

- Based on histogram of local image gradient directions

④ Keypoint descriptor

- Histogram of local gradient directions – Vector with $8 \times (4 \times 4) = 128$ dimensions
- Vector normalization